

[Get started](#)[Open in app](#)[Follow](#)

569K Followers



Photo credit: Pexels

Multi Class Text Classification With Deep Learning Using BERT

Natural Language Processing, NLP, Hugging Face

[Get started](#)[Open in app](#)

Most of the researchers submit their research papers to academic conference because its a faster way of making the results available. Finding and selecting a suitable conference has always been challenging especially for young researchers.

However, based on the previous conferences proceeding data, the researchers can increase their chances of paper acceptance and publication. We will try to solve this text classification problem with deep learning using [BERT](#).

Almost all the code were taken from this [tutorial](#), the only difference is the data.

The Data

The [dataset](#) contains 2,507 research paper titles, and have been manually classified into 5 categories (i.e. conferences) that can be downloaded from [here](#).

Explore and Preprocess

```
1 import torch
2 from tqdm.notebook import tqdm
3
4 from transformers import BertTokenizer
5 from torch.utils.data import TensorDataset
6
7 from transformers import BertForSequenceClassification
8
9 df = pd.read_csv('data/title_conference.csv')
10 df.head()
```

conf_explore.py hosted with ❤ by GitHub

[view raw](#)

conf_explore.py

	Title	Conference
0	Innovation in Database Management: Computer Sc...	VLDB
1	High performance prime field multiplication fo...	ISCAS
2	enchanted scissors: a scissor interface for su...	SIGGRAPH

[Get started](#)[Open in app](#)

Table 1

```
df['Conference'].value_counts()
```

```
ISCAS      864
INFOCOM    515
VLDB       423
WWW        379
SIGGRAPH   326
Name: Conference, dtype: int64
```

Figure 1

You may have noticed that our classes are imbalanced, and we will address this later on.

Encoding the Labels

```
1 possible_labels = df.Conference.unique()
2
3 label_dict = {}
4 for index, possible_label in enumerate(possible_labels):
5     label_dict[possible_label] = index
6 label_dict
```

label_encoding.py hosted with ❤ by GitHub

[view raw](#)

label_encoding.py

```
{'VLDB': 0, 'ISCAS': 1, 'SIGGRAPH': 2, 'INFOCOM': 3, 'WWW': 4}
```

```
df['label'] = df.Conference.replace(label_dict)
```

[Get started](#)[Open in app](#)

Train and Validation Split

Because the labels are imbalanced, we split the data set in a stratified fashion, using this as the class labels.

Our labels distribution will look like this after the split.

```
1  from sklearn.model_selection import train_test_split
2
3  X_train, X_val, y_train, y_val = train_test_split(df.index.values,
4                                                  df.label.values,
5                                                  test_size=0.15,
6                                                  random_state=42,
7                                                  stratify=df.label.values)
8
9  df['data_type'] = ['not_set']*df.shape[0]
10
11 df.loc[X_train, 'data_type'] = 'train'
12 df.loc[X_val, 'data_type'] = 'val'
13
14 df.groupby(['Conference', 'label', 'data_type']).count()
```

train_test_split.py hosted with ❤ by GitHub

[view raw](#)

train_test_split.py

			Title
Conference	label	data_type	
INFOCOM	3	train	438
		val	77
ISCAS	1	train	734
		val	130
SIGGRAPH	2	train	277
		val	49
		train	359



Get started

Open in app

WWW	4	train	322
		val	57

Figure 2

BertTokenizer and Encoding the Data

Tokenization is a process to take raw texts and split into tokens, which are numeric data to represent words.

- Constructs a BERT tokenizer. Based on WordPiece.
- Instantiate a pre-trained BERT model configuration to encode our data.
- To convert all the titles from text into encoded form, we use a function called `batch_encode_plus`, and we will proceed train and validation data separately.
- The 1st parameter inside the above function is the title text.
- `add_special_tokens=True` means the sequences will be encoded with the special tokens relative to their model.
- When batching sequences together, we set `return_attention_mask=True`, so it will return the attention mask according to the specific tokenizer defined by the `max_length` attribute.
- We also want to pad all the titles to certain maximum length.
- We actually do not need to set `max_length=256`, but just to play it safe.
- `return_tensors='pt'` to return PyTorch.
- And then we need to split the data into `input_ids`, `attention_masks` and `labels`.
- Finally, after we get encoded data set, we can create training data and validation data.

```
1 tokenizer = BertTokenizer.from_pretrained('bert-base-uncased',
```

[Get started](#)[Open in app](#)

```
5     df[df.data_type=='train'].Title.values,
6     add_special_tokens=True,
7     return_attention_mask=True,
8     pad_to_max_length=True,
9     max_length=256,
10    return_tensors='pt'
11 )
12
13 encoded_data_val = tokenizer.batch_encode_plus(
14     df[df.data_type=='val'].Title.values,
15     add_special_tokens=True,
16     return_attention_mask=True,
17     pad_to_max_length=True,
18     max_length=256,
19     return_tensors='pt'
20 )
21
22
23 input_ids_train = encoded_data_train['input_ids']
24 attention_masks_train = encoded_data_train['attention_mask']
25 labels_train = torch.tensor(df[df.data_type=='train'].label.values)
26
27 input_ids_val = encoded_data_val['input_ids']
28 attention_masks_val = encoded_data_val['attention_mask']
29 labels_val = torch.tensor(df[df.data_type=='val'].label.values)
30
31 dataset_train = TensorDataset(input_ids_train, attention_masks_train, labels_train)
32 dataset_val = TensorDataset(input_ids_val, attention_masks_val, labels_val)
```

tokenizer_encoding.py hosted with ❤ by GitHub

[view raw](#)

tokenizer_encoding.py

BERT Pre-trained Model

We are treating each title as its unique sequence, so one sequence will be classified to one of the five labels (i.e. conferences).

- bert-base-uncased is a smaller pre-trained model.
- Using num_labels to indicate the number of output labels.

[Get started](#)[Open in app](#)

We also don't need `output_hidden_states`.

```
1 model = BertForSequenceClassification.from_pretrained("bert-base-uncased",
2                                                     num_labels=len(label_dict),
3                                                     output_attentions=False,
4                                                     output_hidden_states=False)
```

BERT_pretrained_model.py hosted with ❤ by GitHub

[view raw](#)

BERT_pretrained_model.py

Data Loaders

- `DataLoader` combines a dataset and a sampler, and provides an iterable over the given dataset.
- We use `RandomSampler` for training and `SequentialSampler` for validation.
- Given the limited memory in my environment, I set `batch_size=3`.

```
1 from torch.utils.data import DataLoader, RandomSampler, SequentialSampler
2
3 batch_size = 3
4
5 dataloader_train = DataLoader(dataset_train,
6                               sampler=RandomSampler(dataset_train),
7                               batch_size=batch_size)
8
9 dataloader_validation = DataLoader(dataset_val,
10                                   sampler=SequentialSampler(dataset_val),
11                                   batch_size=batch_size)
```

data_loaders.py hosted with ❤ by GitHub

[view raw](#)

data_loaders.py

Optimizer & Scheduler

- To construct an optimizer, we have to give it an iterable containing the parameters to optimize. Then, we can specify optimizer-specific options such as the learning rate, epsilon, etc.

[Get started](#)[Open in app](#)

Create a scheduler with a learning rate that decreases linearly from the initial learning rate set in the optimizer to 0, after a warmup period during which it increases linearly from 0 to the initial learning rate set in the optimizer.

```
1  from transformers import AdamW, get_linear_schedule_with_warmup
2
3  optimizer = AdamW(model.parameters(),
4                    lr=1e-5,
5                    eps=1e-8)
6
7  epochs = 5
8
9  scheduler = get_linear_schedule_with_warmup(optimizer,
10                                             num_warmup_steps=0,
11                                             num_training_steps=len(data_loader_train)*epochs)
```

optimizer_scheduler.py hosted with ❤ by GitHub

[view raw](#)

optimizer_scheduler.py

Performance Metrics

We will use f1 score and accuracy per class as performance metrics.

```
1  from sklearn.metrics import f1_score
2
3  def f1_score_func(preds, labels):
4      preds_flat = np.argmax(preds, axis=1).flatten()
5      labels_flat = labels.flatten()
6      return f1_score(labels_flat, preds_flat, average='weighted')
7
8  def accuracy_per_class(preds, labels):
9      label_dict_inverse = {v: k for k, v in label_dict.items()}
10
11     preds_flat = np.argmax(preds, axis=1).flatten()
12     labels_flat = labels.flatten()
13
14     for label in np.unique(labels_flat):
15         y_preds = preds_flat[labels_flat==label]
16         y_true = labels_flat[labels_flat==label]
17         print(f'Class: {label_dict_inverse[label]}')
```


[Get started](#)[Open in app](#)

performance_metrics.py

Training Loop

```
1  import random
2
3  seed_val = 17
4  random.seed(seed_val)
5  np.random.seed(seed_val)
6  torch.manual_seed(seed_val)
7  torch.cuda.manual_seed_all(seed_val)
8
9  def evaluate(dataloader_val):
10
11     model.eval()
12
13     loss_val_total = 0
14     predictions, true_vals = [], []
15
16     for batch in dataloader_val:
17
18         batch = tuple(b.to(device) for b in batch)
19
20         inputs = {'input_ids':      batch[0],
21                  'attention_mask': batch[1],
22                  'labels':         batch[2],
23                  }
24
25         with torch.no_grad():
26             outputs = model(**inputs)
27
28             loss = outputs[0]
29             logits = outputs[1]
30             loss_val_total += loss.item()
31
32             logits = logits.detach().cpu().numpy()
33             label_ids = inputs['labels'].cpu().numpy()
34             predictions.append(logits)
35             true_vals.append(label_ids)
36
```

[Get started](#)[Open in app](#)

```
40     true_vals = np.concatenate(true_vals, axis=0)
41
42     return loss_val_avg, predictions, true_vals
43
44 for epoch in tqdm(range(1, epochs+1)):
45
46     model.train()
47
48     loss_train_total = 0
49
50     progress_bar = tqdm(dataloader_train, desc='Epoch {:1d}'.format(epoch), leave=False,
51 for batch in progress_bar:
52
53     model.zero_grad()
54
55     batch = tuple(b.to(device) for b in batch)
56
57     inputs = {'input_ids':      batch[0],
58               'attention_mask': batch[1],
59               'labels':        batch[2],
60               }
61
62     outputs = model(**inputs)
63
64     loss = outputs[0]
65     loss_train_total += loss.item()
66     loss.backward()
67
68     torch.nn.utils.clip_grad_norm_(model.parameters(), 1.0)
69
70     optimizer.step()
71     scheduler.step()
72
73     progress_bar.set_postfix({'training_loss': '{:.3f}'.format(loss.item()/len(batch
74
75
76     torch.save(model.state_dict(), f'data_volume/finetuned_BERT_epoch_{epoch}.model')
77
78     tqdm.write(f'\nEpoch {epoch}')
79
80     loss_train_avg = loss_train_total/len(dataloader_train)
81     tqdm.write(f'Training loss: {loss train avg}')
```

[Get started](#)[Open in app](#)

```
84     val_f1 = f1_score_func(predictions, true_vals)
85     tqdm.write(f'Validation loss: {val_loss}')
86     tqdm.write(f'F1 Score (Weighted): {val_f1}')
```

training_loop.py hosted with ❤ by GitHub

[view raw](#)

training_loop.py

Epoch 1

Training loss: 0.9007002512753849

Validation loss: 0.6143069127574563

F1 Score (Weighted): 0.7791319217695921

Epoch 2

Training loss: 0.5381144283001613

Validation loss: 0.6438471145765294

F1 Score (Weighted): 0.8207824902152685

Epoch 3

Training loss: 0.35893184876292417

Validation loss: 0.723008230609435

F1 Score (Weighted): 0.8463474188661483

Epoch 4

Training loss: 0.2692523200199349

Validation loss: 0.7796335518272365

F1 Score (Weighted): 0.8341132163207956

Epoch 5

Training loss: 0.18156354463565766

Validation loss: 0.8108082735081321

F1 Score (Weighted): 0.8441012614273822

Figure 3

[Get started](#)[Open in app](#)

```
1 model = BertForSequenceClassification.from_pretrained("bert-base-uncased",
2                                                     num_labels=len(label_dict),
3                                                     output_attentions=False,
4                                                     output_hidden_states=False)
5
6 model.to(device)
7
8 model.load_state_dict(torch.load('data_volume/finetuned_BERT_epoch_1.model', map_location=device))
9
10 _, predictions, true_vals = evaluate(dataloader_validation)
11 accuracy_per_class(predictions, true_vals)
```

loading_evaluating.py hosted with ❤ by GitHub

[view raw](#)

loading_evaluating.py

Class: VLDB
Accuracy: 45/64

Class: ISCAS
Accuracy: 124/130

Class: SIGGRAPH
Accuracy: 29/49

Class: INFOCOM
Accuracy: 65/77

Class: WWW
Accuracy: 33/57

Figure 4

[Jupyter notebook](#) can be found on [Github](#). Enjoy the rest of the weekend!

[Get started](#)[Open in app](#)

By Towards Data Science

Every Thursday, the Variable delivers the very best of Towards Data Science: from hands-on tutorials and cutting-edge research to original features you don't want to miss. [Take a look.](#)

[Get this newsletter](#)[NLP](#)[Nlp Tutorial](#)[Text Classification](#)[Document Classification](#)[Machine Learning](#)[About](#) [Write](#) [Help](#) [Legal](#)

Get the Medium app

